

## First we read the data just as told

```
In [ ]: from sklearn.datasets import fetch_lfw_people
lfw_people = fetch_lfw_people(min_faces_per_person=150, resize=0.4)
X = lfw_people.data
Y = lfw_people.target
lfw_people.target_names
```

```
Out[ ]: array(['Colin Powell', 'George W Bush'], dtype='<U13')
```

## Then we take a look the data

```
In [ ]: lfw_people.images.shape
```

```
Out[ ]: (766, 50, 37)
```

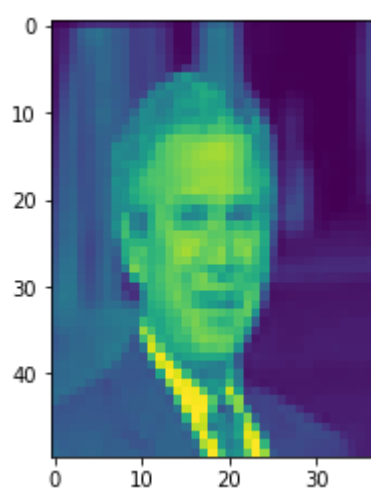
```
In [ ]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.25, shuffle=True)
print(f"x_train shape: {x_train.shape}")
print(f"x_test shape: {x_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

```
x_train shape: (574, 1850)
x_test shape: (192, 1850)
y_train shape: (574,)
y_test shape: (192,)
```

## Let's see some examples of our dataset:

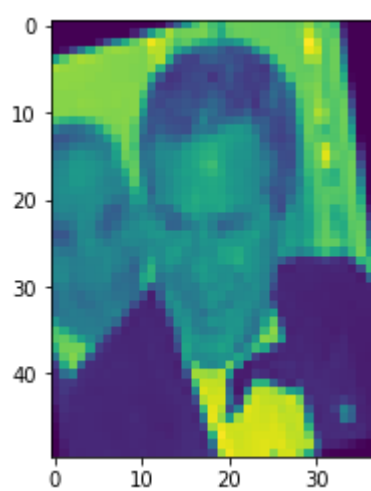
```
In [ ]: import matplotlib.pyplot as plt
plt.imshow(x_train[0].reshape(50, 37))
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x7f1849b324c0>
```



```
In [ ]: plt.imshow(x_test[0].reshape(50, 37))
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x7f18499b47f0>
```



## The actual data looks like:

```
In [ ]: x_train
```

```
Out[ ]: array([[0.01699346, 0.0248366 , 0.05751634, ..., 0.1267974 , 0.09281046,
0.08627451],
[0.00130719, 0.          , 0.          , ..., 0.          , 0.          ,
0.          ],
[0.00392157, 0.00392157, 0.0130719 , ..., 0.00915033, 0.          ,
0.          ],
...,
[0.          , 0.          , 0.          , ..., 0.075817 , 0.075817 ,
0.07189543],
[0.3934641 , 0.4993464 , 0.5751634 , ..., 0.10849673, 0.15816994,
0.23006536],
[0.          , 0.          , 0.          , ..., 0.          , 0.          ,
0.          ]], dtype=float32)
```

For a better result; first we normalize the data:

Pay attention that we first fit over x\_train then transform over x\_train and x\_test

```
In [ ]: from sklearn.preprocessing import StandardScaler
scalar = StandardScaler().fit(x_train)
x_train = scalar.transform(x_train)
x_test = scalar.transform(x_test)
```

Here we define a function that does ALL the stuff we need!!!

It gets pca\_n\_components as the argument and feed it to PCA. Then use it to reduce the number of features. Prints some info. then fit the model over train and test and prints the results and scores. It also plots confusion matrices.

```
In [ ]: from sklearn.decomposition import PCA
def DoATest(pca_n_components, model):
    global x_train, x_test, y_train, y_test
    pca = PCA(n_components=pca_n_components).fit(x_train)
    pca_applied_x_train = pca.transform(x_train)
    pca_applied_x_test = pca.transform(x_test)
    print("Applied PCA:")
    print(f"x_train shape: {pca_applied_x_train.shape}")
    print(f"x_test shape: {pca_applied_x_test.shape}")

    model.fit(pca_applied_x_train, y_train)
    print("TESTING OVER TRAIN PART OF DATA:")
    print("RESULTS OVER Y_TRAIN:")
    report_results(model.predict(pca_applied_x_train), y_train)
    plt.show()

    print("TESTING OVER TEST PART OF DATA:")
    print("RESULTS OVER Y_TEST:")
    report_results(model.predict(pca_applied_x_test), y_test)
    plt.show()

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns


def report_results(y_predicted, y_true, figure_ax=None):
    print(f"accuracy: {accuracy_score(y_true, y_predicted)}")
    print(classification_report(y_true, y_predicted))
    mat = confusion_matrix(y_true, y_predicted)
    sns.heatmap(mat, annot=True, ax=figure_ax)
```

Now all we need to do is to play with the parameters and just call the function we mentioned earlier to find the best parameter set!

```
In [ ]: from sklearn.neural_network import MLPClassifier as MLP
DoATest(200, MLP())
```


```
Applied PCA:
x_train shape: (574, 200)
x_test shape: (192, 200)
TESTING OVER TRAIN PART OF DATA:
RESULTS OVER Y_TRAIN:
accuracy: 1.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:


accuracy:	0.953125				
	precision	recall	f1-score	support	
0	0.93	0.92	0.92	59	
1	0.96	0.97	0.97	133	
accuracy			0.95	192	
macro avg	0.95	0.94	0.94	192	
weighted avg	0.95	0.95	0.95	192	



```
In [ ]: DoATest(150, MLP())
```


Applied PCA:  
x\_train shape: (574, 150)  
x\_test shape: (192, 150)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0

	precision	recall	f1-score	support	
0	1.00	1.00	1.00	177	
1	1.00	1.00	1.00	397	
accuracy			1.00	574	
macro avg	1.00	1.00	1.00	574	
weighted avg	1.00	1.00	1.00	574	



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:


accuracy:	0.953125				
	precision	recall	f1-score	support	
0	0.93	0.92	0.92	59	
1	0.96	0.97	0.97	133	
accuracy			0.95	192	
macro avg	0.95	0.94	0.94	192	
weighted avg	0.95	0.95	0.95	192	



```
In [ ]: DoATest(100, MLP())
```


Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:  
accuracy: 0.9375


	precision	recall	f1-score	support
0	0.91	0.88	0.90	59
1	0.95	0.96	0.96	133
accuracy			0.94	192
macro avg	0.93	0.92	0.93	192
weighted avg	0.94	0.94	0.94	192



```
In [ ]: DoATest(80, MLP())
```


Applied PCA:  
x\_train shape: (574, 80)  
x\_test shape: (192, 80)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:  
accuracy: 0.9322916666666666


	precision	recall	f1-score	support
0	0.91	0.86	0.89	59
1	0.94	0.96	0.95	133
accuracy			0.93	192
macro avg	0.93	0.91	0.92	192
weighted avg	0.93	0.93	0.93	192



```
In [ ]: DoATest(100, MLP(max_iter=500))
```

Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0


	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:

accuracy: 0.9427083333333334


	precision	recall	f1-score	support
0	0.94	0.86	0.90	59
1	0.94	0.98	0.96	133
accuracy			0.94	192
macro avg	0.94	0.92	0.93	192
weighted avg	0.94	0.94	0.94	192



```
In [ ]: DoATest(100, MLP(max_iter=500, hidden_layer_sizes=(150,)))
```

Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0


	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:

accuracy: 0.9375


	precision	recall	f1-score	support
0	0.91	0.88	0.90	59
1	0.95	0.96	0.96	133
accuracy			0.94	192
macro avg	0.93	0.92	0.93	192
weighted avg	0.94	0.94	0.94	192



```
In [ ]: DoATest(100, MLP(max_iter=500, hidden_layer_sizes=(150,), activation="logistic"))
```


Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	177
1	1.00	1.00	1.00	397
accuracy			1.00	574
macro avg	1.00	1.00	1.00	574
weighted avg	1.00	1.00	1.00	574



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:  
accuracy: 0.9479166666666666

	precision	recall	f1-score	support
0	0.92	0.92	0.92	59
1	0.96	0.96	0.96	133
accuracy			0.95	192
macro avg	0.94	0.94	0.94	192
weighted avg	0.95	0.95	0.95	192




```
In [ ]: DoATest(100, MLP(max_iter=500, hidden_layer_sizes=(150,)), activation="logistic", solver="sgd"))
```

Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:  
accuracy: 0.9895470383275261

	precision	recall	f1-score	support
0	1.00	0.97	0.98	177
1	0.99	1.00	0.99	397
accuracy			0.99	574
macro avg	0.99	0.98	0.99	574
weighted avg	0.99	0.99	0.99	574


/home/pedram/.local/lib/python3.8/site-packages/sklearn/neural\_network/\_multilayer\_perceptron.py:702: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (500) reached and the optimization hasn't converged yet.  
warnings.warn(



TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:

accuracy: 0.9583333333333334

	precision	recall	f1-score	support
0	0.95	0.92	0.93	59
1	0.96	0.98	0.97	133
accuracy			0.96	192
macro avg	0.96	0.95	0.95	192
weighted avg	0.96	0.96	0.96	192




```
In [ ]: DoATest(100, MLP(max_iter=700, hidden_layer_sizes=(150,), activation="logistic", solver="sgd"))
```

Applied PCA:  
x\_train shape: (574, 100)  
x\_test shape: (192, 100)  
TESTING OVER TRAIN PART OF DATA:  
RESULTS OVER Y\_TRAIN:

accuracy: 0.9947735191637631

	precision	recall	f1-score	support
0	1.00	0.98	0.99	177
1	0.99	1.00	1.00	397
accuracy			0.99	574
macro avg	1.00	0.99	0.99	574
weighted avg	0.99	0.99	0.99	574

/home/pedram/.local/lib/python3.8/site-packages/sklearn/neural\_network/\_multilayer\_perceptron.py:702: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (700) reached and the optimization hasn't converged yet.  
warnings.warn(




TESTING OVER TEST PART OF DATA:  
RESULTS OVER Y\_TEST:

accuracy: 0.96875

	precision	recall	f1-score	support
0	0.96	0.93	0.95	59
1	0.97	0.98	0.98	133
accuracy			0.97	192
macro avg	0.97	0.96	0.96	192
weighted avg	0.97	0.97	0.97	192





Perfect!

The last one is the best indeed! Because at the same time that has a good accuracy(96.87%) on the test data, it has very close accuracy on the train(99.4) which shows us that the model is accurate and ALSO it hasn't overfit! because the accuracy over train and test are pretty close!

(The earlier tries that the accuracy over train were 100% and accuracy over test was much different the model had overfit)